

Rapid Acquisition of Disjunction from Prosody and Consistency

Masoud Jasbi

masoudj@stanford.edu
Department of Linguistics
Stanford University

Akshay Jaggi

ajaggi@stanford.edu
Department of Linguistics
Stanford University

Michael C. Frank

mcf Frank@stanford.edu
Department of Psychology
Stanford University

Abstract

Ambiguity poses a challenge to children learning the meaning of words. For example, a disjunction such as *A or B* can be interpreted as inclusive (*A or B*, or both) or exclusive (*A or B*, not both). Linguistic studies suggest that the core meaning of *or* is inclusive and despite the dominance of exclusive interpretations, children successfully and quickly learn the meaning of *or* as inclusive. This raises a learning puzzle: how can children quickly learn what they rarely hear? We argue that children can use the regularities of *or* in child-directed speech to learn the interpretation of a disjunction from a few examples.

Keywords: language acquisition; word learning; logical words; and; or; nativism; constructivism.

Introduction

The social media company LADbible reported the following in a tweet: “James Bond producer says next 007 could be black **or** a woman”. A twitter user named Robert responded sarcastically with: “if only women could be black!” What in the producer’s speech gave the impression that the next 007 could not be both black **and** a woman? The word *or*. A disjunction like “*A or B*” is associated with two interpretations: **inclusive**, and **exclusive**. “*A or B*” is inclusive when it is interpreted as “*A or B or both*”. This is probably what LADbible meant when reporting the James Bond producer. However, “*A or B*” can also be exclusive: “*A or B*, but **not both**”. Robert’s response shows that he had an exclusive interpretation of *or*. What factors determine the interpretation of *or* and how do children learn its meaning given this ambiguity?

A large body of research in linguistics and philosophy in the past 50 years has created a common consensus on the meaning of *or* (see Aloni (2016)). Data on the interpretation of disjunction across different sentences, contexts, and even languages show that the core meaning of disjunction words such as *or* is **inclusive**. This is similar to the definition of disjunction in formal logics. The **exclusive** interpretation of *or* is the result of enhancing its inclusive semantics via other (extra-semantic) factors such as intonation (Pruitt & Roelofsen, 2013), inconsistency of the options (Geurts, 2006), and pragmatic reasoning over the speaker’s choice of *or* instead of *and* (Grice, 1975). Therefore, interpreting a disjunction is a complex process that needs to take into account the meaning of *or* as well as different structural and contextual factors that accompany it. How do we, as children, learn such a complex interpretive system?

There are two accounts of children’s acquisition of disjunction: 1. a **constructivist** account, and 2. a **nativist** account. Under the constructivist account, children learn the meaning of *or* by paying attention to how parents use it in

different contexts. They form usage-rules and expand their usage repertoire as they grow up. The prediction is that children’s production of *or* is slow and gradual, mirroring what they hear from parents. Morris (2008) found that *and* is about 13 times more frequent than *or* in parents’ speech to children. As predicted by the constructivist account, Morris (2008) reported that children also learn to produce *and* much more quickly than *or*. They reach the adult rate of *and* at age 3 while for *or* there is a gradual increase in production, possibly reaching the adult level at ages 5 or 6. The faster acquisition of *and* is consistent with the constructivist theory that emphasizes the role of usage frequency in children’s linguistic development. Morris (2008) also reported that the majority of *or* examples children hear are exclusive. He argued that consistent with the constructivist account, the majority of *or*’s children produce are also exclusive. Therefore, the inclusive semantics of *or* is developed after the exclusive interpretation.

However, several comprehension studies of *or* in different linguistic contexts show that children as young as three-years old interpret *or* as inclusive disjunction (Crain, 2012). This is surprising given Morris (2008)’s finding that the majority of *or* examples children hear are exclusive. How do children learn the interpretation of disjunction as inclusive if they rarely hear it? We call this the **learning puzzle of disjunction**. Crain (2012) suggests that instead of learning from the parents’ usage of *or*, children rely on the innate knowledge that the disjunction operator in their native language must have an inclusive meaning. This nativist account predicts that *or* is learned relatively quickly and accurately by children.

Here we present an alternative answer that provides a synthesis between the nativist and constructivist accounts. We show that children can use regularities in parents’ usage of *or* to learn the interpretation of disjunction from a handful of examples. In study 1, we use a large scale corpus study of parents and children’s speech to show that both *and* and *or* appear relatively quickly in children’s speech and reach the adult rate by the age 4. This finding is consistent with the comprehension studies that show children have an adult-like understanding of these words by the age four. In study 2, we replicate Morris (2008)’s finding that the majority of *or* examples in child-directed speech have an exclusive interpretation. However, we also show that these exclusive interpretations correlate systematically with two factors external to *or*: intonation and consistency of the options. Exclusive interpretations are either inconsistent in nature (e.g. clean or dirty) or carry a distinct rise-fall intonation. We show that setting aside these cases, the interpretation of a disjunction is inclusive.

We argue that if children track the interpretive cues that

accompany a disjunction, they can tease apart the role of the word *or* from factors that accompany and enhance it to shape the exclusive interpretation. This way children can discover that exclusivity correlates with rise-fall intonation and inconsistent options, while *or* itself does not exclude the option of both disjuncts being true. We implement this novel account in a decision-tree that correctly learns to predict the interpretation of a disjunction with 80% accuracy after only a few examples. Our results show that the richness and systematicity of children’s linguistic input allows rapid acquisition of disjunction with no need for an innate assumptions specific to the meaning of disjunction. We discuss the important implications of this work for the theories of word learning in the last section.

Study 1: Corpus Study

First, we conducted an exploratory and large scale investigation of *and* and *or* productions in parents and children. The goal of the study was to find out when children start producing these words and when they reach the adult rate of production. We conclude that children start producing *and* around 1.5 or 2 years of age, and *or* between the ages of 2 and 3. They reach the adult rate of production for *and* around 3 and for *or* around 4 or possibly earlier.

Methods

We accessed the Child Language Data Exchange System (CHILDES, MacWhinney (2000)) via the online platform *childes-db* and its associated R package *childesr* (Sanchez et al., in prep). We extracted all instances of *and* and *or* from the English corpora (ENG-NA and ENG-UK). We limited our analysis to the data between one and six years because there is scarce data outside this age range. We computed the relative frequency of connective production by dividing the total number of *and/or* in the speech of fathers, mothers, and children at a particular age by the total number of words spoken at that age. We present the relative frequency as parts per thousand.

Results

In figure 1, we show the relative frequencies of *and* and *or* in the speech of parents and children between one and six years. It is important to note that the y-axes for *and* vs. *or* show different ranges. This is due to the big difference in the relative frequencies of *and* and *or*. In the speech of parents, *and* is produced around 20 times per thousand words while *or* is only produced around 2 times per thousand words. This confirms previous findings that *or* is much less frequent in child directed speech than a similar function word such as *and*.




Figure 1: The relative frequency of AND (top) and OR (bottom) per thousand words in the speech of fathers, mothers, and children between the ages of 1 and 6.

And and *or* seem to show different developmental trajectories in the speech of children. For *and*, there is a rapid increase in its production between the ages of 1.5 and 3 before it reaches the adult rate around the age 3 and stay at that level until the age 6. For *or*, on the other hand, we see a slow increase from the age 2 until the age 6 when it reaches the adult rate. This difference in the development of *and* & *or* production was attributed to the frequency of these items in child-directed speech. Since *and* is much more frequent than *or*, it is learned much faster than *or*. Morris (2008) argued that such patterns support the item-based and usage-based acquisition of logical words.

However, the analysis above does not control for other factors that can affect the production of words by children. An important factor to control for is the development of speech acts. While content words such as *dog* or *chair* may appear freely in different types of speech acts, function words are highly constrained by the type of speech acts they can appear in. For example, it is reasonable to assume that question words such as *how* and *why* are much more likely to occur in questions than statements (declaratives). If parent-child interaction is such that parents ask more questions than children, it is not surprising to find higher rates of *how* and *why* production in parents than children. Therefore, it is important to control for the speech act a function word appears in.

Figure 2 shows the relative frequencies of *and* and *or* in questions vs. declaratives, in the speech of parents and children between one and six years. Here, the relative frequency is computed by dividing the total number of *and/or* in a question/declarative in the speech of fathers, mothers, and children at a particular age, by the total number of words in a question/declarative spoken at that age. As before, we present the relative frequency as parts per thousand.

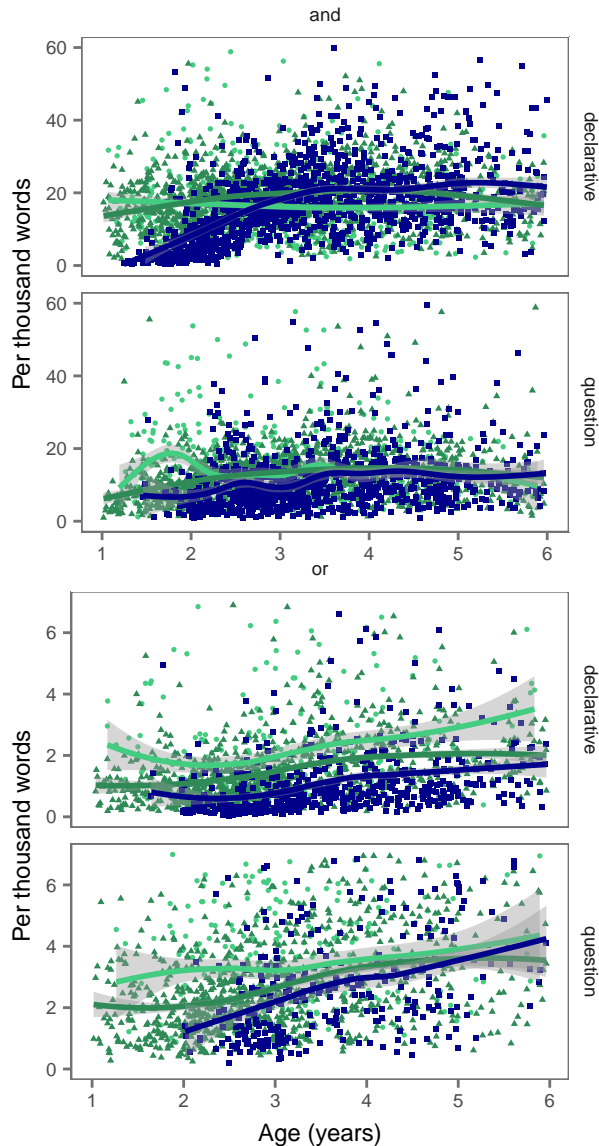


Figure 2: The relative frequency of AND (top) and OR (bottom) per thousand words in declaratives and questions in the speech of fathers, mothers, and children between the ages of 1 and 6,

The results show similar developmental trajectories for the production of *and* and *or* in children. For both words, there is a relatively rapid increase in their frequency between the ages of 2 and 4 before reaching the parent rate at the age of 4 and staying at that rate until the age of 6. This pattern of

production is consistent with the nativist observations that the acquisition of *and* and *or* is rapid and that children have an adult-like comprehension of these two connectives at the age 4.

Summary

And is a lot more frequent than or in child directed speech.

In the first six years, it appears that children reach the adult production rate for *and* but not *or*.

This is at least partly because *or* is more frequent in questions and children produce fewer questions than parents.

The developmental trajectory of connective production is best described as a quick increase in production between the ages of 2 and 4 and staying around the parents' rate between the ages 4 and 6. This is compatible with the comprehension studies which suggest children understand *and* and *or* by the age four.

Study 2: Annotation Study

Second, we conducted an detailed and small-scale investigation of *or* productions in parents and children. The goal of the study was to find out how children learn the meaning of *or* from such little data. We conclude that children hear significantly more exclusive *or* supporting Morris (2008)'s findings. Additionally, we find that the exclusive uses of *or* are rich in structures that index the exclusive interpretation. These indicative structures include specifically intonation and consistency of the disjuncts. Exclusive interpretations are either inconsistent in nature (e.g. clean or dirty) or carry a distinct rise-fall intonation. We then show that a decision tree model can learn to accurately (>80%) predict exclusivity using this information after seeing few (<20) examples.

Methods

We accessed the Providence corpus via CHILDES (Demuth, Culbertson, & Alter, 2006). We extracted all instances of *or* along with the two utterances before and after the utterance containing *or*. We annotated the examples for four major categories. *Exclusivity Interpretation*: This category represents the goal task: understanding the intended form of disjunction. *Intonation*: This category was divided into three main intonation patterns over the disjuncts. Intonation could be flat, rising overall, or rise-fall. These three categories were selected because rise-fall in particular has been shown to lead to an exclusive interpretation of the disjunction (???). *Consistency*: This category tracked the logical consistency of the disjuncts. Two disjuncts were marked as inconsistent if replacing the word “or” with “and” produced a logical conflict. For example in “Are your feet clean or dirty?”, the disjuncts are inconsistent: the addressee's feet cannot be both clean and dirty.

Two raters annotated the same 240 utterances to develop a reliability score. The interrater reliability was calculated over 8 iterations of 30 examples. Training only completed after 3 consecutive iterations with reliability over 0.7 for all categories.




Figure 3: Distribution of exclusive and inclusive interpretations broken down by intonation and consistency. Error bars represent bootstrapped 95% confidence intervals

Results

First, similar to Morris (2008), we found that the majority of *or* examples in CDS receive an exclusive interpretation (~65). Figure 3 shows this difference in distribution. However, the rate of exclusive interpretations change systematically when we break the data down by prosody and consistency (figure below). A mixed-effects binomial logistic regression with the fixed effects of intonation, consistency, and random effects for children found intonation and consistency significant in interpreting disjunctions. Disjunctions were more likely to be interpreted as exclusive if they received a rise-fall intonation ($\beta=-3.79$, $z=1.66$, $p < 0.001$) or if they were inconsistent ($\beta=-2.2$, $z=2.08$, $p < 0.001$). Disjunctions were more likely to be interpreted as inclusive if they were consistent and received a rising intonation ($\beta=0.58$, $z=0.24$, $p < 0.001$) or flat intonation ($\beta=0.38$, $z=0.27$, $p < 0.001$).

Figure 4 demonstrates the clear relationship between exclusive interpretation, rise-fall intonation, and inconsistent disjuncts. Without these markers of exclusivity, the inclusive interpretations remain. This supports Crain (2012)’s claim that the “default” interpretation of “or” is inclusive.

Next, using Sci-kit Learn’s Decision Tree Module (Pedregosa et al., 2011), we built a predictive model to train on annotated *or* utterances and predict the exclusivity of unseen *or* utterances (annotated for intonation and consistency). Averaged over 100 trials and training on 200 examples, the average accuracy of a binary tree was 83%. More remarkably, the tree achieved an average of 80% accuracy after training on only 20 examples. The success of such a simple tree indicates that children could use a simple model to rapidly learn the exclusive interpretation of *or* from little data.

Summary

In study 2, we confirm Morris (2008)’s finding that exclusive interpretations of *or* are far more common than inclusive interpretations. However, we also show that the majority of these exclusive interpretations coincide with systematic indicators. Disjunctions that are accompanied by rise-fall intonation or inconsistent disjuncts are far more likely to be exclusive. Disjunctions that do not bear these features are more likely to be inclusive. Accounting for these external factors, a

simple decision tree can rapidly learn to predict the exclusive interpretation of the disjunction.

Conclusion

List what you showed and argued for.

Talk about important implications for theories of word learning.

Possibly talk about the original tweet and how that falls still outside what your algorithm learns.

Acknowledgements

Place acknowledgments (including funding information) in a section at the end of the paper.

References

- Aloni, M. (2016). Disjunction. In E. N. Zalta (Ed.), *The stanford encyclopedia of philosophy*. Stanford University. Retrieved from <https://plato.stanford.edu/archives/win2016/entries/disjunction/>
- Crain, S. (2012). *The emergence of meaning*. Cambridge University Press.
- Demuth, K., Culbertson, J., & Alter, J. (2006). Word-minimality, epenthesis and coda licensing in the early acquisition of english. *Language and Speech*, 49(2), 137–173.
- Geurts, B. (2006). Exclusive disjunction without implicatures. *Ms., University of Nijmegen*.
- Grice, H. P. (1975). Logic and conversation. In P. Cole & J. Morgan (Eds.), *Syntax and semantics* (Vol. 3: Speech Acts, pp. 43–58). Academic Press.
- MacWhinney, B. (2000). *The childe project: The database* (Vol. 2). Psychology Press.
- Morris, B. J. (2008). Logically speaking: Evidence for item-based acquisition of the connectives and & or. *Journal of Cognition and Development*, 9(1), 67–88.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... others. (2011). Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12(Oct), 2825–2830.
- Pruitt, K., & Roelofsen, F. (2013). The interpretation of prosody in disjunctive questions. *Linguistic Inquiry*, 44(4), 632–650.
- Sanchez, A., Meylan, S., Braginsky, M., MacDonald, K., Yurovsky, D., & Frank, M. C. (in prep). Childes-db: A flexible and reproducible interface to the child language data exchange system (childes).




Figure 4: Distribution of exclusive and inclusive interpretations broken down by intonation and consistency. Error bars represent bootstrapped 95% confidence intervals




Figure 5: Decision Tree Accuracy as a function of number of examples seen